

Train an SSD network in a self-driving car application

1. Import Required Libraries

Load essential libraries like `torch`, `torchvision`, `PIL`, `cv2`, `numpy`, and `requests` for deep learning, image processing, and I/O.

2. Define COCO Class Labels

Specify class names from the COCO dataset, which the SSD model uses for labeling detected objects.

3. Load Pretrained SSD300 VGG16 Model

Load the pretrained `ssd300_vgg16` model from `torchvision` and switch it to evaluation mode (`model.eval()`).

4. Prepare Image Transformations

Define a transformation pipeline using `torchvision.transforms` to resize and convert the image to tensor format expected by the model.

5. Load Image from URL or Local Path

Attempt to load the image using `requests` from a URL. If it fails, fallback to a local image path using `PIL.Image.open`.

6. Preprocess Image for SSD Input

Resize the image to 300x300, convert it to a tensor, and add a batch dimension before passing it to the model.

7. Run Inference with the Model

Use `torch.no_grad()` for inference and get predictions like bounding boxes, labels, and confidence scores.

8. Filter and Rescale Detections

Filter predictions by confidence threshold (>0.5) and rescale bounding boxes from model input size to original image dimensions.

9. Draw Bounding Boxes and Labels

For each detected object (e.g., car, person), draw a bounding box and add the label with confidence score on the original image using `cv2`.

10. Display the Result

Convert the image back to RGB and display it using `matplotlib.pyplot` without axes.

