

Building a simple Generative Adversarial Network (GAN) using TensorFlow

1. Set Hyperparameters and Random Seed

The random seed is set for reproducibility. Hyperparameters like latent dimension, epochs, batch size, buffer size, and image shape (28x28 grayscale) are defined.

2. Load and Preprocess the MNIST Dataset

The MNIST dataset is loaded and reshaped to match the input requirements of the GAN. The images are normalized to a range of $[-1, 1]$ for better training stability, and the dataset is batched and shuffled.

3. Define the Generator Model

The generator network is built using a series of dense layers, LeakyReLU activations, Batch Normalization layers, and a final dense layer that reshapes the output into the MNIST image shape. The output activation function is `tanh` to scale pixel values between -1 and 1.

4. Define the Discriminator Model

The discriminator is built as a binary classifier. It takes the image as input, flattens it, and passes it through a few dense layers with LeakyReLU activations, and ends with a sigmoid activation to predict whether the image is real or fake.

5. Loss Functions

Discriminator Loss: Binary cross-entropy loss is computed for both real and fake images. The goal of the discriminator is to classify real images as 1 and fake images as 0.

Generator Loss: The generator's goal is to fool the discriminator, so its loss is calculated as binary cross-entropy with a target of 1 (since the generator wants to make fake images appear real).

6. Initialize Models and Optimizers

The generator and discriminator are initialized using the defined models. Adam optimizers are used for both networks with a learning rate of `1e-4`.

7. Training Step Function

In each training step:

- Noise is generated as the input to the generator.
- The generator creates fake images from the noise.
- The discriminator evaluates both real and fake images.
- Losses for both generator and discriminator are calculated, and gradients are updated using backpropagation.

8. Generate and Save Images

At every 10th epoch, images are generated using a fixed input (`seed`) to visualize how the generator is learning over time. These images are saved as PNG files.

9. Training Loop

The GAN is trained for the specified number of epochs (50 in this case). For each batch of images from the dataset, the `train_step` function is executed. After every 10 epochs, generated images are saved.

10. Final Image Generation and Model Saving

After training, the generator produces final images, and the model is saved to a file (`generator_model.h5`). The final generated images are also displayed.